

From "What Broke" to "What Changed"

Michael Beemer, Dynatrace

Parth Suthar, DevCycle



hosted at: <https://beeme1mr.github.io/kubecon-us-25-flag-observability/>



KubeCon



CloudNativeCon

North America 2025



Michael Beemer
Senior Product Manager

 [beeme1mr](#)



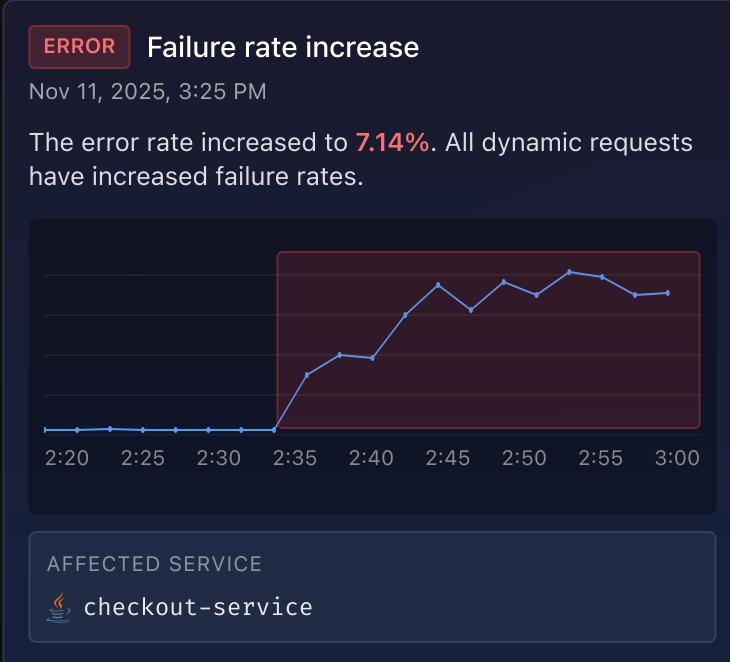
Parth Suthar
Software Engineer

 [suthar26](#)

3:25 PM Monday

A familiar story...

- **ERROR** 3:25 PM
Failure rate increase detected
- **INVESTIGATION** 3:30 PM
No recent deployments found
- **INVESTIGATION** 4:05 PM
No code changes in last 48h
- **INVESTIGATION** 4:50 PM
All dependencies healthy
- **ESCALATION** 5:20 PM
All hands on deck
- **ROOT CAUSE** 6:25 PM
"Oh, I toggled that feature flag"



⚠️ The Missing Context
Feature flag changes are **invisible** to traditional monitoring



Feature flags are **hidden** from observability tools
making it **difficult** to pinpoint changes
as the **root cause** of incidents

The Observability Gap

What traditional monitoring misses

⌚ What We See

- ⚠ Error rates spike
- ⌚ Latency increases
- ✗ Failed requests
- ⚡ Resource exhaustion

Traditional metrics & traces

∅ What's Hidden

- 🚩 Which flag changed?
- ⌚ When was it toggled?
- 👤 Who was affected?
- 💡 What variant was served?

Feature flag context

A Real-World Example: OpenTelemetry Demo

Feature flags that trigger production-like issues

Astronomy Shop

- Full microservices e-commerce platform
- 10+ services, multiple languages
- Production-grade observability stack
- Uses **OpenFeature** for feature flagging

🎯 Perfect Test Bed

Includes flags that deliberately enable problems to demonstrate observability challenges

Problem Scenarios

⚠️ `recommendationServiceCacheFailure`
Memory leak → 1.4x exponential growth → OOM crashes

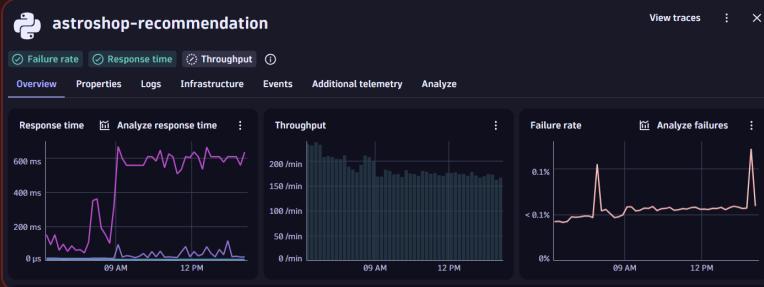
⚠️ `paymentServiceUnreachable`
Invalid endpoint → all payment requests fail

⚠️ `imageSlowLoad`
Fault injection → 5s image load latency

Impact vs Root Cause

Observability shows the **symptoms**, but hides the **diagnosis**

What Your Dashboard Shows



- ⚠️ Clear Impact: Response time spiked from 145ms to 600ms!
- ❓ Unknown Cause: No code changes, no deploys...

What's Actually Happening



The Hidden Truth

- ✅ Quick fix: Toggle flag off → problem gone
- ⌚ Gives time to debug properly
- 🎯 But monitoring tools can't see this connection

The Mitigation Problem:

Hours spent restarting pods, rolling back code, and debugging — when a **30-second flag toggle** would have stopped the bleeding

Why This Matters

The hidden cost of invisible feature flags



Slower Recovery

MTTR increases when the easiest fix is invisible



Wrong Direction

Debugging code that isn't actually broken



Escalation

Simple toggle becomes all-hands incident



We need feature flags as **first-class concept** in observability

Feature Flag Observability

A progressive approach for teams looking to gain visibility into flags



Level 0

Flying Blind



Level 1

Broadcast Blast



Level 2

Manual Events



Level 3

Auto Mapping



Level 4

Trace-Level

Level 0: Flying Blind



You see something is wrong, but have no idea why

The Problem

- ⌚ No visibility into flag changes
- ⚠ Pure guesswork during incidents
- ⌚ Hours of manual hunting

Checkout Service

What You See

⚠ Failure rate spiking!

But what changed? No deployments, no code changes...

Time to check logs, metrics, Slack, coffee machine...

⚠ The Reality

A feature flag was toggled 30 seconds ago, but you have no way to know that.

Level 1: Broadcast Blast



Same flag change annotated on all services... but only one is actually using the flag

Checkout service

Payment service

Recommendation service

⚠ The Red Herring Problem

When you manually configure events, you might send them to services that **don't use the flag**. This creates noise during incidents and can send investigations in the wrong direction.

"We see the flag changed at the same time, but these services look fine... maybe it's not the flag?"

Level 2: Manual Change Events 🤝

Send events to specific services... but requires manual mapping and can become outdated

checkout-service

payment-service

recommendation-service

⚠ The Manual Mapping Problem

You can **manually configure** which services receive flag change events, but this requires maintaining a mapping. As your system evolves, these mappings become **outdated** — events might go to services that no longer use the flag, or miss new services that started using it.

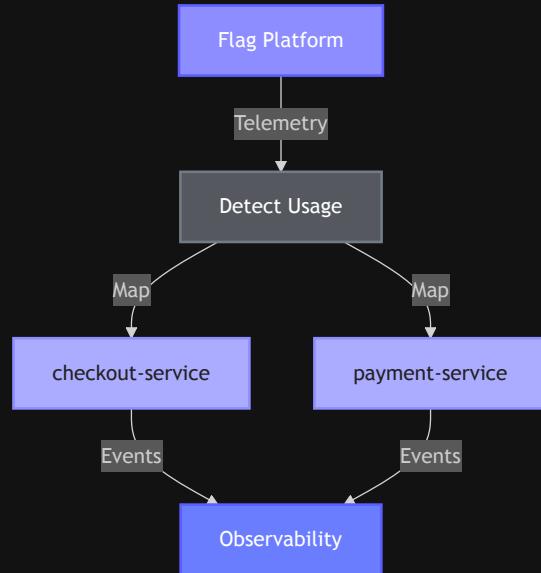
"We configured this 6 months ago... did we update it when we refactored?"

Level 3: Automatic Event Mapping



Telemetry-driven routing eliminates manual configuration

Intelligent Routing



Benefits

- ✓ **Automatic discovery**
Detects where flags are evaluated
- ✓ **Smart routing**
Events only to affected services
- ✓ **Zero maintenance**
Stays current as code evolves
- ⚠ **Correlation ≠ Causation**
Good indicator, but can't measure exact impact

Solves the scaling problem, but requires a more complicated setup. How can it be improved further?

Level 4: Trace-Level Observability

Feature flags as a first-class observability concept

Fine-Grained Insights

Response Time

-  **See flag evaluations**
Every trace shows which flags were evaluated
-  **Realtime mapping of flag presence per service**
Filter by traces with specific flag presence
-  **Filter by flag presence**
Isolate requests that used a specific flag
-  **Compare variants**
See performance differences in real-time
-  **Instant root cause**
Immediately identify which variant caused issues

 Split by variant - "on" variant shows 3200-3800ms! Root cause identified.



Trace-level observability makes
feature flags a **first-class citizens**
in your monitoring stack



Feature Flag Observability Standards

Bridging feature flags and observability through open standards

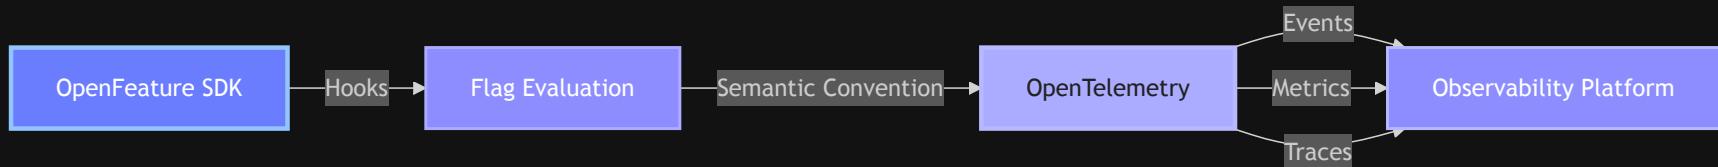
OpenFeature 🤝 OpenTelemetry

OpenFeature

- Open standard for **feature flagging**
- Vendor-neutral SDK
- Hook system** for extensibility

OpenTelemetry

- Open standard for **observability**
- Metrics, traces, and logs
- Semantic conventions** for consistency



Feature Flag Semantic Convention

Standardized attributes that make flag evaluations observable

Obvious Attributes

`feature_flag.key`

The flag identifier (e.g., "new-checkout")

`feature_flag.result.variant`

Which variant was served (e.g., "on", "treatment")

`feature_flag.result.value`

The value returned by the flag (e.g., true, false, "blue")

`feature_flag.provider.name`

Flag provider (e.g., "flagd", "launchdarkly", "devcycle")

`feature_flag.result.reason`

Why this variant? (e.g., "targeting_match", "default")

Nuanced Attributes

`feature_flag.set.id`

Human-readable logical identifier for where this flag is managed
"acme-org/web-app/production"

`feature_flag.context.id`

Provider's context identifier (fallback: targeting key)
Used by providers like DevCycle to lookup user evaluations and simulation

`feature_flag.version`

Ruleset version at evaluation time
Track which configuration was active (number, hash, etc.)

How It Works: OpenFeature Hooks

```
import { MyProvider } from 'my-flag-provider';
import { EventHook } from '@openfeature/open-telemetry-hooks';

// Initialize OpenFeature with your flag provider
await OpenFeature.setProviderAndWait(new MyProvider());

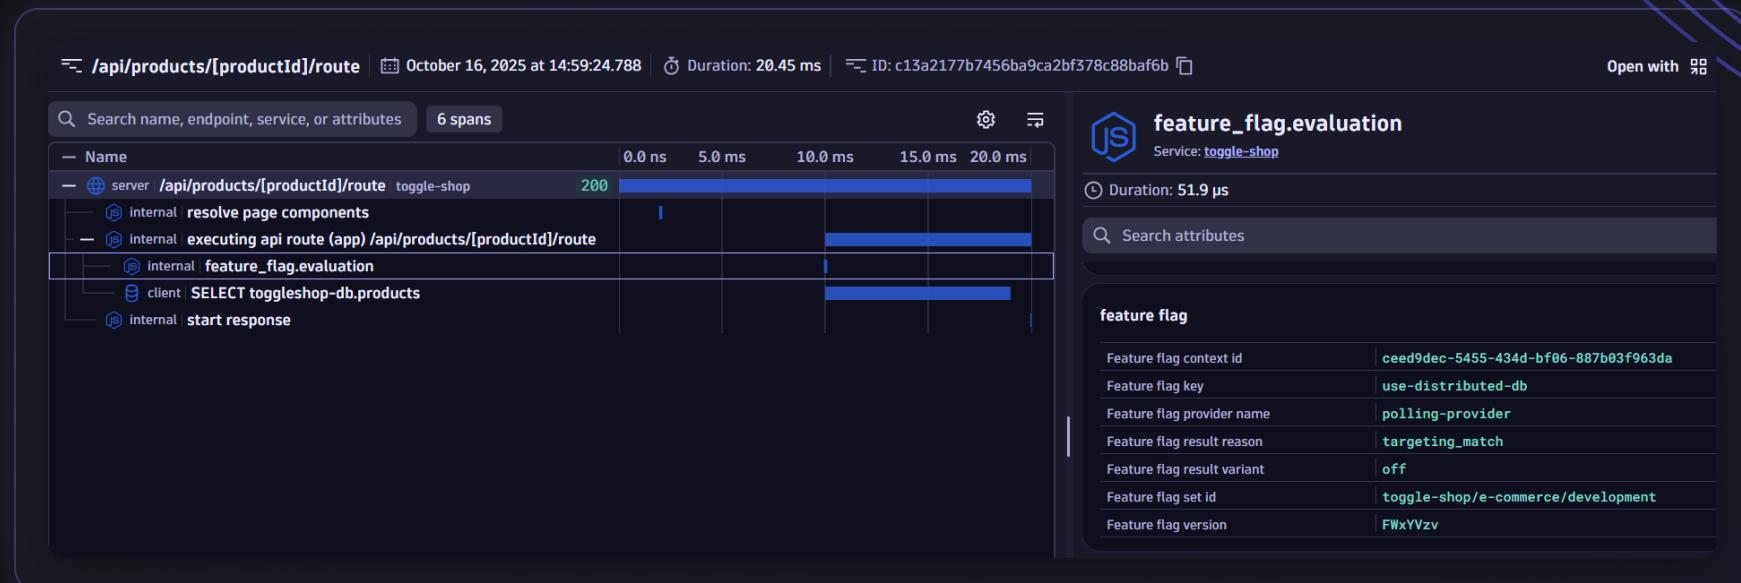
// Register the OpenTelemetry hook
OpenFeature.addHooks(new EventHook());
const client = OpenFeature.getClient();

// Prepare evaluation context
const context = {
  targetingKey: 'user_123',
  tier: 'premium',
  environment: 'production'
};

// Evaluate flag - automatically traced!
const variant = await client.getStringValue(
  'new-checkout-flow',
  'control',
  context
);
```

Seeing It in Action

Feature flag evaluation captured in distributed tracing



Why Standards Matter



Interoperability

Works across vendors and tools



Consistency

Same attributes everywhere



Adoption

Easy to implement and adopt





Progressive Delivery with Observability

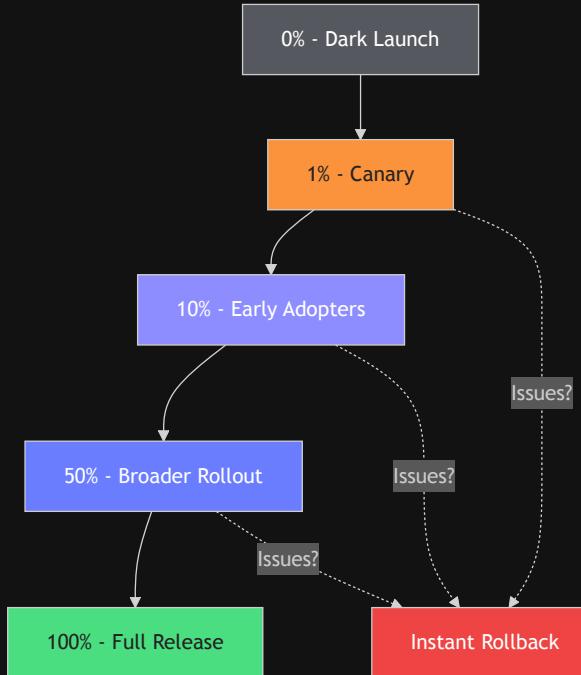
Transforming theory into practice

What Is Progressive Delivery?

Concepts

- Gradually expose customers to new features
- Monitor impact at each step
- Roll back instantly if problems arise
- Minimize risk, maximize confidence

Rollout Stages



Rollout targeting setup

Targeting Rules

— Hide

1 Name Gradual rollout

Definition All Users Comparator Enter a value and press enter...

+ Add Definition

Serve Variation On

Schedule Multi-Step Rollout 1% Current Rollout %

1 1% on Nov 4, 2025 12:00 AM EST (GMT-5)

2 10% on Nov 6, 2025 12:00 AM EST (GMT-5)

3 50% on Nov 11, 2025 12:00 AM EST (GMT-5)

4 100% on Nov 13, 2025 12:00 AM EST (GMT-5)

+ Add Rollout Step

Randomize Using User ID

↑ ↓

The screenshot shows a 'Targeting Rules' configuration interface. At the top, there's a 'Name' field containing 'Gradual rollout'. Below it, a 'Definition' section includes a dropdown for 'All Users' and a 'Comparator' dropdown with an input field 'Enter a value and press enter...'. A '+ Add Definition' button is also present. The 'Serve' section has a dropdown set to 'Variation On'. The 'Schedule' section is set to 'Multi-Step Rollout' and shows a current rollout percentage of '1%'. Below this, four rollout steps are defined: 1) 1% on Nov 4, 2025 12:00 AM EST (GMT-5); 2) 10% on Nov 6, 2025 12:00 AM EST (GMT-5); 3) 50% on Nov 11, 2025 12:00 AM EST (GMT-5); and 4) 100% on Nov 13, 2025 12:00 AM EST (GMT-5). A '+ Add Rollout Step' button is available. The 'Randomize Using' section is set to 'User ID'. At the bottom right, there are up and down arrow buttons.

Observability Enables Progressive Delivery

Variant-level visibility reveals the root cause



new-checkout-flow @ checkout-service



Rolled Back

0% rollout (100% control / 0% treatment)

Traffic Split

Error Rate

P95 Latency



Rollback Triggered - Treatment Variant Disabled

Without Flag Observability:

Hours debugging aggregate metrics, unclear which change caused the issue

With Flag Observability:

Instant variant identification & rollback in seconds

Key Benefits

✖ Without Observability

- ⌚ Blind rollouts
- ⌚ Slow to detect issues
- ⌚ Manual rollback decisions
- ⌚ Can't compare variants
- ⚠ High-risk deployments

✓ With Observability

- ⌚ Data-driven rollouts
- ⚡ Instant issue detection
- ⌚ Automated rollback triggers
- ⌚ Real-time variant comparison
- ✓ Low-risk deployments

Observability is the **key** to progressive delivery

Challenges and Opportunities

The journey isn't complete

⚠ Current Limitations

⌚ Scaling Challenges

Users must handle deduplication, sampling, and aggregation strategies themselves

🔌 Provider Dependency

Requires OpenFeature providers to include semantic convention attributes

🚀 Opportunities Ahead

🔗 Metrics Convention

Standardize flag evaluation metrics for consistent monitoring

⭐ Change Events

Define standard format for flag configuration changes

嚚 Flags as Entities

Model feature flags as first-class entities in observability platforms

Let's Build the Future Together

Feature flags are powerful tools for modern software delivery

But they need to be visible in our observability stack

OpenFeature + OpenTelemetry make this possible with open standards

Join us in transforming how we understand and deploy software

From "What Broke" to "What Changed"

Thank You!

Questions?

Find us after the talk or online:



Resources

- ↗ openfeature.dev
- ⌚ github.com/open-openfeature/
- ↖ opentelemetry.io



Connect

- ♫ CNCF Slack: #openfeature
- ♫ CNCF Slack: #opentelemetry
- 🐦 @openfeature / @opentelemetry